

START

mizan (ميزان) — the balance

An AI that answers Qur'an & hadith questions — then proves every quote.

☐ VERIFIED · ☐ REJECTED · ? UNVERIFIABLE — computed, never guessed

\$ bun run demo → VERIFIED + REJECTED in ~1s, offline, no key

27,234 real records · Qur'an + 4 Sunan + Muwatta · Track 04

THE PROBLEM

AI makes up religious text — and it looks real

Fluent words are not proof.

AI hallucinates fluent hadith that does not exist

Fake and real quotes look identical to readers

Fuzzy checkers fail: fake-but-plausible hadith $\approx$ 97% (CWE-345)

In worship & belief, a wrong quote is not a typo

THE SOLUTION

Answer first — then prove every quote

Three badges. One rule: exact containment earns VERIFIED.

VERIFIED → quote literally inside the cited source

REJECTED → source exists, quote NOT in it (+ real nearby text shown)

UNVERIFIABLE → no citation / paraphrase / timeout (reason named)

Every badge ships with: quoted text + source row + source URL

Six small checks — one path to VERIFIED

First failure decides the badge.

1 Quote? · 2 Citation? · 3 ≤3 cites? · 4 Record found?

5 Inside it? exact folded match → VERIFIED · 6 Proof present?

Folds tashkeel/tatweel/alef/digits — NEVER changes a letter

True quote (with or without tashkeel) passes · one-word fake never does

Three doors — same verifier

CLI, web page, MCP server.

A · Terminal: `bun run demo` (offline, ~1s) · `bun run ask "..."`

B · Web: `bun run demo-server` → `localhost:3000`

C · MCP: `bun run mcp` → read-only tool: `verify`

Output per claim: `verified` / `rejected` / `unverifiable`

Real input → real output

About one second in the terminal.

IN: bun run ask "oneness of God?"

OUT: [✅ VERIFIED] ikhlas-1 · quoted: قل هو الله أحد · quran 6222

OUT: [✅ REJECTED] one-word-changed hadith + real nearby text

OUT: [? UNVERIFIABLE] paraphrase — reason named

Rules no one can bypass

Machine-checked gates, not promises.

Exact containment only — no similarity, no percentages

Single VERIFIED site; CI fails if a second appears (G-6)

Grades stored as dataset says, or null — never invented

Logs carry hashes, never your question

What we do NOT claim

An over-claim here is a religious error.

No Bukhari / Muslim yet — biggest gap, stated upfront

Paraphrase □ UNVERIFIABLE (26 human-ruled cases)

Default = labelled replay; live needs your API key

Proves the QUOTE exists — not the interpretation

Grade = dataset's grade, never our ruling

REMEMBER

Problem · Solution · Workflow

نعلم الآلة .. لتخدم الرسالة — with proof.

Problem: AI invents hadith that look real

Solution: check each quote — exact match only

Workflow: ask → answer → computed badge + evidence

Try: bun install && bun run demo